

**COSC76/276 Artificial Intelligence**  
**Fall 2022**  
**Adversarial search**


Soroush Vosoughi  
Computer Science  
Dartmouth College  
Soroush@Dartmouth.edu

# Logistics

- PA-2 due Oct 10th at 11:59pm ET
- Assignments are being steadily graded and made available
- 3 extra late days has been given

# Recap: Partial observations

- Belief states: where the agent thinks it might be



- Upper bound of belief state space size:  $2^3 - 1$

# Recap: Adversarial search and minimax


- Adversarial search as a tree search
- Minimax algorithm to find optimal strategy against optimal opponent

# Minimax

## Example

```
def value(state):  
    if the state is a terminal state: return the state's utility  
    if the next agent is MAX: return max-value(state)  
    if the next agent is MIN: return min-value(state)
```


```
def max-value(state):  
    initialize v =  $-\infty$   
    for each successor of state:  
        v = max(v, value(successor))  
    return v
```



```
def min-value(state):  
    initialize v =  $+\infty$   
    for each successor of state:  
        v = min(v, value(successor))  
    return v
```

Max

Min



# Today's learning objectives

- Implement solutions that can practically work for real-world adversarial search

# Outline

- Alpha-beta pruning
- Handling resource limits


# Outline

- Alpha-beta pruning
- Handling resource limits




# Alpha-beta pruning intuition

- Consider a node  $n$  somewhere in the tree such that Player has a choice of moving to that node




# Alpha-beta pruning intuition

- Consider a node  $n$  somewhere in the tree such that Player has a choice of moving to that node
  - If Player has a better choice  $m$  either at the parent node of  $n$  or at any choice point further up, then  $n$  will never be reached in actual play



# Alpha-beta pruning intuition example



# Alpha beta pruning

- Alpha: MAX's best option on path to root
- Beta: MIN's best option on path to root

**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
     $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
    **return** the *action* in ACTIONS(*state*) with value *v*

**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
     $v \leftarrow -\infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
         $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
        **if**  $v \geq \beta$  **then return** *v*  
         $\alpha \leftarrow \text{MAX}(\alpha, v)$   
    **return** *v*

**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
     $v \leftarrow +\infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
         $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
        **if**  $v \leq \alpha$  **then return** *v*  
         $\beta \leftarrow \text{MIN}(\beta, v)$   
    **return** *v*

# Example


**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
**return** the *action* in  $\text{ACTIONS}(\text{state})$  with value  $v$

**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if**  $\text{TERMINAL-TEST}(\text{state})$  **then return**  $\text{UTILITY}(\text{state})$   
 $v \leftarrow -\infty$   
**for each**  $a$  **in**  $\text{ACTIONS}(\text{state})$  **do**  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    **if**  $v \geq \beta$  **then return**  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
**return**  $v$

**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if**  $\text{TERMINAL-TEST}(\text{state})$  **then return**  $\text{UTILITY}(\text{state})$   
 $v \leftarrow +\infty$   
**for each**  $a$  **in**  $\text{ACTIONS}(\text{state})$  **do**  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    **if**  $v \leq \alpha$  **then return**  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
**return**  $v$

Max

Min



Let's work through  
this example together  
with a full run of the  
algorithm

# Alpha-beta pruning example following the code

**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
     $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
    **return** the *action* in  $\text{ACTIONS}(\text{state})$  with value  $v$


**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
     $v \leftarrow -\infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
         $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
        **if**  $v \geq \beta$  **then return**  $v$   
         $\alpha \leftarrow \text{MAX}(\alpha, v)$   
    **return**  $v$

**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
     $v \leftarrow +\infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
         $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
        **if**  $v \leq \alpha$  **then return**  $v$   
         $\beta \leftarrow \text{MIN}(\beta, v)$   
    **return**  $v$

Alpha: MAX's best option on path to root  
Beta: MIN's best option on path to root

Max

Min



# Alpha-beta pruning example following the code

**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
**return** the *action* in  $\text{ACTIONS}(\text{state})$  with value  $v$


**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if**  $\text{TERMINAL-TEST}(\text{state})$  **then return**  $\text{UTILITY}(\text{state})$   
 $v \leftarrow -\infty$   
**for each**  $a$  **in**  $\text{ACTIONS}(\text{state})$  **do**  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    **if**  $v \geq \beta$  **then return**  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
**return**  $v$

**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if**  $\text{TERMINAL-TEST}(\text{state})$  **then return**  $\text{UTILITY}(\text{state})$   
 $v \leftarrow +\infty$   
**for each**  $a$  **in**  $\text{ACTIONS}(\text{state})$  **do**  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    **if**  $v \leq \alpha$  **then return**  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
**return**  $v$

Alpha: MAX's best option on path to root  
Beta: MIN's best option on path to root

Max

Min



# Alpha-beta pruning example following the code

**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
**return** the *action* in  $\text{ACTIONS}(\text{state})$  with value  $v$


**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if**  $\text{TERMINAL-TEST}(\text{state})$  **then return**  $\text{UTILITY}(\text{state})$   
 $v \leftarrow -\infty$   
**for each**  $a$  **in**  $\text{ACTIONS}(\text{state})$  **do**  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    **if**  $v \geq \beta$  **then return**  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
**return**  $v$

**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if**  $\text{TERMINAL-TEST}(\text{state})$  **then return**  $\text{UTILITY}(\text{state})$   
 $v \leftarrow +\infty$   
**for each**  $a$  **in**  $\text{ACTIONS}(\text{state})$  **do**  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    **if**  $v \leq \alpha$  **then return**  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
**return**  $v$

Alpha: MAX's best option on path to root  
Beta: MIN's best option on path to root

Max

Min





# Alpha-beta pruning example following the code

function ALPHA-BETA-SEARCH(*state*) returns an action  
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
 return the *action* in  $\text{ACTIONS}(\text{state})$  with value  $v$


function MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) returns a utility value  
 if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
 $v \leftarrow -\infty$   
 for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
 if  $v \geq \beta$  then return  $v$   
 $\alpha \leftarrow \text{MAX}(\alpha, v)$   
 return  $v$

function MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) returns a utility value  
 if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
 $v \leftarrow +\infty$   
 for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
 if  $v \leq \alpha$  then return  $v$   
 $\beta \leftarrow \text{MIN}(\beta, v)$   
 return  $v$

Alpha: MAX's best option on path to root  
 Beta: MIN's best option on path to root

Max

Min



# Alpha-beta pruning example following the code

**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
**return** the *action* in  $\text{ACTIONS}(\text{state})$  with value  $v$


**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if**  $\text{TERMINAL-TEST}(\text{state})$  **then** **return**  $\text{UTILITY}(\text{state})$   
 $v \leftarrow -\infty$   
**for each**  $a$  **in**  $\text{ACTIONS}(\text{state})$  **do**  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s,a), \alpha, \beta))$   
    **if**  $v \geq \beta$  **then** **return**  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
**return**  $v$

**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if**  $\text{TERMINAL-TEST}(\text{state})$  **then** **return**  $\text{UTILITY}(\text{state})$   
 $v \leftarrow +\infty$   
**for each**  $a$  **in**  $\text{ACTIONS}(\text{state})$  **do**  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s,a), \alpha, \beta))$   
    **if**  $v \leq \alpha$  **then** **return**  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
**return**  $v$

Alpha: MAX's best option on path to root  
Beta: MIN's best option on path to root

Max

Min



# Alpha-beta pruning example following the code

**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
**return** the *action* in  $\text{ACTIONS}(\text{state})$  with value  $v$

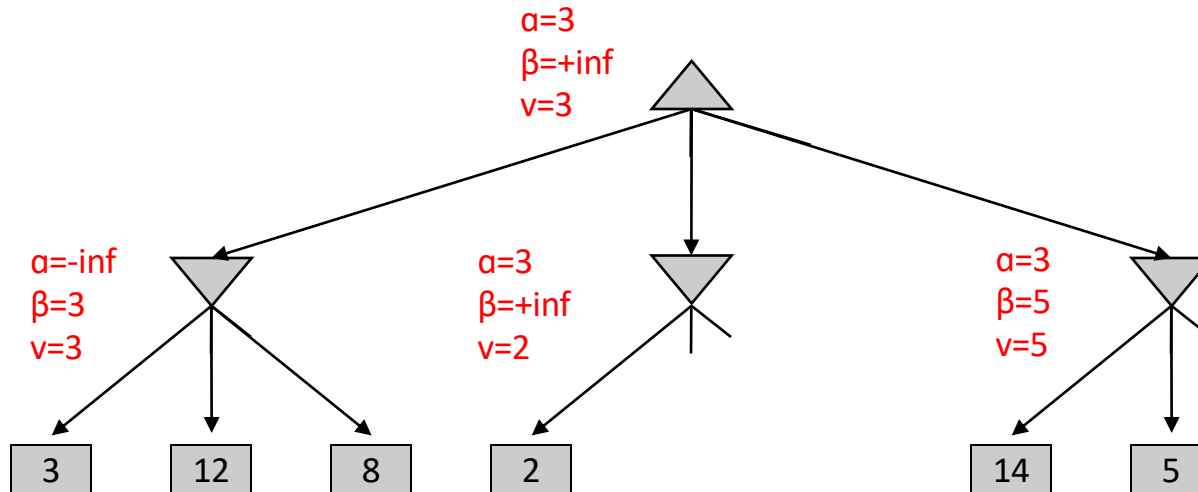
**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if** **TERMINAL-TEST**(*state*) **then return** **UTILITY**(*state*)  
 $v \leftarrow -\infty$   
**for each** *a* **in**  $\text{ACTIONS}(\text{state})$  **do**  
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
**if**  $v \geq \beta$  **then return**  $v$   
 $\alpha \leftarrow \text{MAX}(\alpha, v)$   
**return**  $v$

**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if** **TERMINAL-TEST**(*state*) **then return** **UTILITY**(*state*)  
 $v \leftarrow +\infty$   
**for each** *a* **in**  $\text{ACTIONS}(\text{state})$  **do**  
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
**if**  $v \leq \alpha$  **then return**  $v$   
 $\beta \leftarrow \text{MIN}(\beta, v)$   
**return**  $v$

Alpha: MAX's best option on path to root  
Beta: MIN's best option on path to root

Max

Min



# Alpha-beta pruning example following the code

**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
**return** the *action* in ACTIONS(*state*) with value *v*


Alpha: MAX's best option on path to root  
 Beta: MIN's best option on path to root

**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
 $v \leftarrow -\infty$   
**for each** *a* **in** ACTIONS(*state*) **do**  
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
**if**  $v \geq \beta$  **then return** *v*  
 $\alpha \leftarrow \text{MAX}(\alpha, v)$   
**return** *v*


**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
 $v \leftarrow +\infty$   
**for each** *a* **in** ACTIONS(*state*) **do**  
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
**if**  $v \leq \alpha$  **then return** *v*  
 $\beta \leftarrow \text{MIN}(\beta, v)$   
**return** *v*

Max

Min




# Alpha-beta pruning example




Another view is: what are the possible values that can appear in a branch?

For example, here, it could go from  $-\text{inf}$  to 2

# Alpha-beta pruning example




# Alpha-beta pruning example



Order matters!

What if we'd  
gotten a 2  
before a 14?


# Alpha-Beta Quiz



Anything that  
can be  
pruned?




# Alpha-Beta Quiz



Anything that  
can be  
pruned?


# Alpha-Beta Quiz 2

Anything that  
can be  
pruned?



# Alpha-Beta Quiz 2

Anything that  
can be  
pruned?



# Alpha-beta pruning Properties


- This pruning has **no effect** on minimax value computed for the root
- Good child ordering improves effectiveness of pruning
- With “perfect ordering” Time complexity drops to  $O(b^{m/2})$

# Outline

- Alpha-beta pruning
- Handling resource limits

# Resource limits

- Problem
  - In realistic games, cannot search to leaves, as decisions need to happen in **real-time**
- Solution
  - Cut-off test instead of terminal test
    - E.g., depth-limited search, iterative deepening
  - Replace terminal utilities with **an evaluation function** for non-terminal positions to estimate desirability of position
- Consequence
  - Guarantee of optimal play is gone




# Cut-off search

- In the recursive call, we need to have bookkeeping of the depth

$$\text{H-MINIMAX}(s, d) = \begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN}. \end{cases}$$

# Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



- Ideal function: returns the actual minimax value of the position
- In practice, typically weighted linear sum of features:


$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.  $f_1(s) = (\text{num white queens} - \text{num black queens})$ , etc.



# When to have the cut-off?

- Iterative deepening within allocated time
- **Horizon effects** are a problem: stalling tactics can push bad states beyond the depth searched




# When to have the cut-off?

- Iterative deepening within allocated time
- **Horizon effects** are a problem: stalling tactics can push bad states beyond the depth searched
- **Quiescent states**, i.e., states that will not have a large change in value, are good candidates for cutoff
  - Quiescence search: search “volatile” positions to a greater depth than “stable” ones

# Repeated states

- In games, repeated states occur frequently because of transpositions – i.e., different permutations of the move sequence end up in the same position
  - e.g., [a1, b1, a2, b2] vs. [a1, b2, a2, b1]
- It's worthwhile to store the evaluation of this position in a hash table – **transposition table** – the first time it is encountered
  - similar to the “explored set” in graph-search
- Tradeoff:
  - Transposition table can be too big
  - Which to keep and which to discard

# PA-3: chess game





# Summary


- Alpha-beta pruning
  - Alpha: MAX's best option on path to root
  - Beta: MIN's best option on path to root
  - Prune part of the tree that won't be played given the current values of alpha and beta
- Real-time decisions
  - Terminal-test -> Cut-off test
  - Evaluation function to estimate desirability of a position

# Types of Games

- Classified over different axes:
  - Number of players
  - Zero sum
  - **Deterministic** or **stochastic**
  - **Perfect** or imperfect information




# Stochastic games



In stochastic games, uncertain outcomes controlled by chance, not an adversary (e.g., dice, card shuffling, unpredictable opponents)


# Why not minimax?



Discussion




# Why not minimax?



- Worst case reasoning is too conservative
- Need average case reasoning

# Expectiminimax Search

- In stochastic games, values should reflect average-case (expectiminimax) outcomes, not worst-case (minimax) outcomes
- **Expectiminimax search:** compute the average score under optimal play
  - Max and min nodes as in minimax search
  - Chance nodes are like another player with “actions” with associated probabilities
  - Calculate their **expected utilities**, i.e. weighted average (expectation) of children



# Expectiminimax

EXPECTIMINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

if the next agent is EXP: return exp-value(state)


def max-value(state):

initialize  $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return  $v$



def min-value(state):

initialize  $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$

return  $v$

def exp-value(state):

initialize  $v = 0$

for each successor of state:

$p = \text{probability}(\text{successor})$


$v += p * \text{value}(\text{successor})$

return  $v$

# Expectimax example of chance node


EXPECTIMINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$




$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$

# Expectimax Example




Example


# Expectimax Example



# Expectimax Example

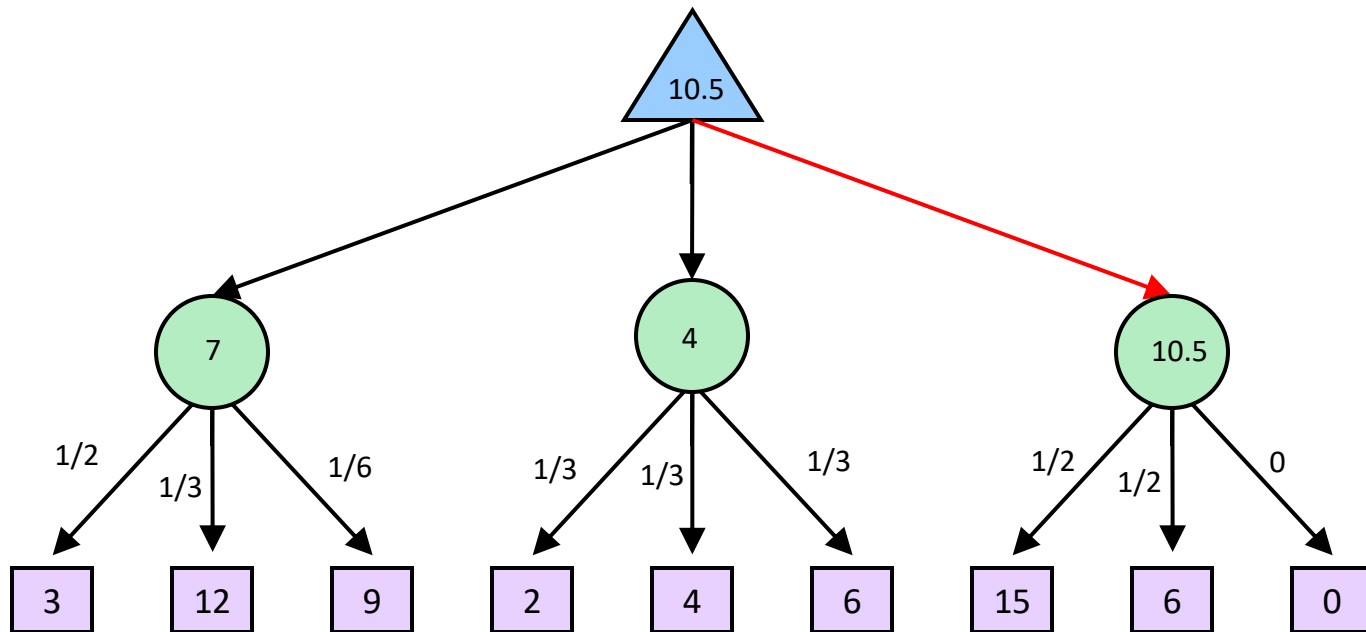


# Expectimax Example





# Expectimax Example





# Model for chance nodes

# Model for chance nodes

- Expectiminimax search assumes to have a probabilistic model of how the opponent (or environment) will behave in any state
  - Model could be a simple uniform distribution (roll a die)
  - Model could be sophisticated and require a great deal of computation
  - We have a chance node for any outcome out of our control: opponent or environment

# Types of Games

- Classified over different axes:
  - Number of players
  - Zero sum
  - Deterministic or stochastic
  - Perfect or imperfect information



# Games with imperfect information

- Games could be partially observable
  - Result into belief states
- Solution for each state in the belief as deterministic game (with all techniques we have seen)
  - Instead of solving all of them, randomly samples states in the belief to solve them

# Commonsense - Information is important

Mondays:

- Road A leads to 1 gold. Road B leads to a fork.
- Go left from the fork, 100 gold. Go right, die.
- *Optimal strategy: B*

Tuesdays:

- Road A leads to 1 gold. Road B leads to a fork.
- Go left from the fork, die. Go right, 100 gold.
- *Optimal strategy: B*

I can't remember what day it is. *Choose B?*

# Commonsense - Information is important

Mondays:

- Road A leads to 1 gold. Road B leads to a fork.
- Go left from the fork, 100 gold. Go right, die.
- *Optimal strategy: B*

Tuesdays:


- Road A leads to 1 gold. Road B leads to a fork.
- Go left from the fork, die. Go right, 100 gold.
- *Optimal strategy: B*

I can't remember what day it is. *Choose B?*

With partial observability, value of an action depends on belief state of the agent -> agent might try to reduce that uncertainty

# Multi-Agent Utilities

- With multiple players, the tree can be extended with additional nodes corresponding to each player
- Generalization of minimax:
  - Terminals have utility tuples
  - Node values are also utility tuples
  - Each player maximizes its own component
  - Can give rise to cooperation and competition dynamically





# Summary

- Stochastic games introduce chance nodes
  - Expectiminimax for finding the solution that maximizes the average case
  - A probabilistic model is needed
- Partially observable games result in belief states

# Next

- Non-zero sum games