

**COSC76/276 Artificial Intelligence**  
**Fall 2022**  
**Adversarial search**

Soroush Vosoughi  
Computer Science  
Dartmouth College  
Soroush@Dartmouth.edu


# Reminders

- PA-2 due today at 11:59pm ET


# Recap: Adversarial search and minimax

- Techniques to make the adversarial search problem more tractable: alpha-beta pruning

# Alpha-beta pruning intuition example




# Alpha-Beta Quiz



Anything that  
can be  
pruned?

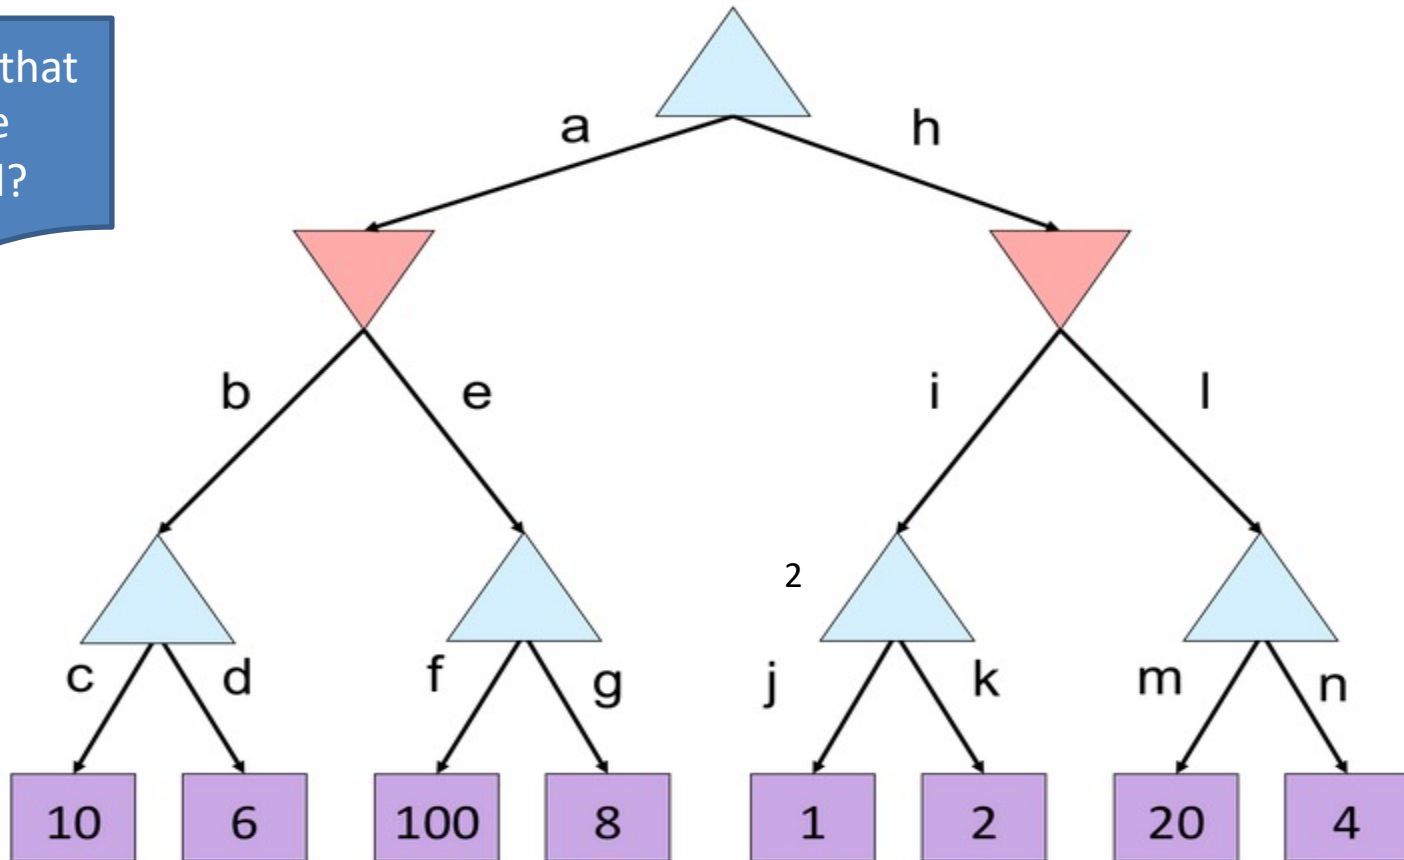
# Alpha-Beta Quiz



Anything that  
can be  
pruned?


# Alpha-Beta Quiz 2

Anything that  
can be  
pruned?



# Alpha-Beta Quiz 2

Anything that  
can be  
pruned?






# Alpha-beta pruning Properties

- This pruning has **no effect** on minimax value computed for the root
- Good child ordering improves effectiveness of pruning
- With “perfect ordering” Time complexity drops to  $O(b^{m/2})$

# Resource limits

- Problem
  - In realistic games, cannot search to leaves, as decisions need to happen **real-time**
- Solution
  - Cut-off test instead of terminal test
    - E.g., depth-limited search, iterative deepening
  - Replace terminal utilities with **an evaluation function** for non-terminal positions to estimate desirability of position
- Consequence
  - Guarantee of optimal play is gone




# Cut-off search

- In the recursive call, we need to have bookkeeping of the depth

$$\begin{aligned} \text{H-MINIMAX}(s, d) = & \\ \left\{ \begin{array}{ll} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN.} \end{array} \right. \end{aligned}$$

# Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



- Ideal function: returns the actual minimax value of the position
- In practice, typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.  $f_1(s) = (\text{num white queens} - \text{num black queens})$ , etc.

# When to have the cut-off?

- **Quiescent states**, i.e., states that will not have a large change in value, are good candidates for cutoff
  - Quiescence search
- **Horizon effects** are a problem: stalling tactics can push bad states beyond the depth searched

# Repeated states

- In games, repeated states occur frequently because of transpositions – i.e., different permutations of the move sequence end up in the same position
  - e.g., [a1, b1, a2, b2] vs. [a1, b2, a2, b1]
- It's worthwhile to store the evaluation of this position in a hash table – **transposition table** – the first time it is encountered
  - similar to the “explored set” in graph-search
- Tradeoff:
  - Transposition table can be too big
  - Which to keep and which to discard

# PA-3: chess game



# Summary

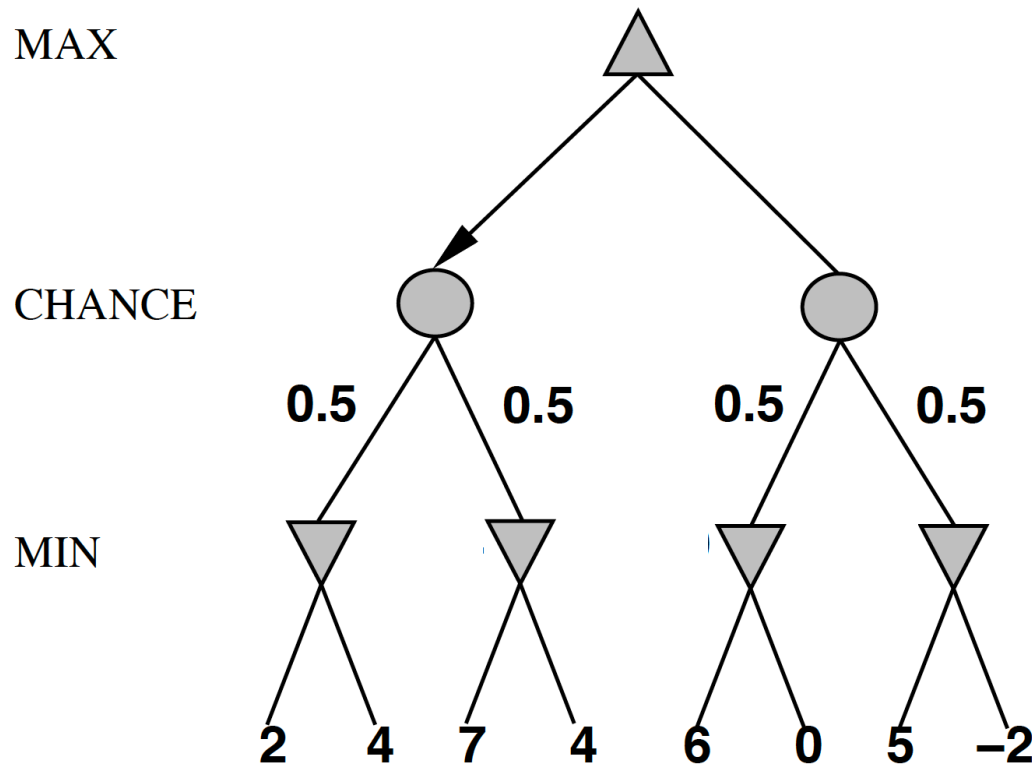
- Alpha-beta pruning
  - Alpha: MAX's best option on path to root
  - Beta: MIN's best option on path to root
  - Prune part of the tree that won't be played given the current values of alpha and beta
- Real-time decisions
  - Terminal-test -> Cut-off test
  - Evaluation function to estimate desirability of a position



# Games


- Classified over different axes:
  - Deterministic or stochastic
  - Perfect or imperfect information
  - Number of players
  - Zero sum

# Stochastic games



In stochastic games, uncertain outcomes controlled by chance, not an adversary (e.g., dice, card shuffling, unpredictable opponents)


# Why not minimax?



- Worst case reasoning is too conservative
- Need average case reasoning

# Expectiminimax Search

- In stochastic games, values should reflect average-case (expectiminimax) outcomes, not worst-case (minimax) outcomes
- **Expectiminimax search:** compute the average score under optimal play
  - Max and min nodes as in minimax search
  - Chance nodes are like another player with “actions” with associated probabilities
  - Calculate their **expected utilities**, i.e. weighted average (expectation) of children



# Expectiminimax

EXPECTIMINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

if the next agent is EXP: return exp-value(state)


def max-value(state):

initialize  $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return  $v$



def min-value(state):

initialize  $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$

return  $v$

def exp-value(state):

initialize  $v = 0$

for each successor of state:

$p = \text{probability}(\text{successor})$


$v += p * \text{value}(\text{successor})$

return  $v$

# Expectimax example of chance node


EXPECTIMINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$




$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$

# Expectimax Example




Example

# Expectimax Example







# Expectimax Example




# Expectimax Example



# Expectimax Example




# Expectimax Pruning?



Can we prune?


Discussion

# Expectimax Pruning?



Can we prune? Last nodes can be worth a lot.

# Expectimax Pruning?



Can we prune? Last nodes can be worth a lot.

What if the values are bounded?

# Similar tricks for real-time decisions

- Alpha-beta pruning
  - Bound for utility values
- Cut-off tests
  - Evaluation function
  - Monte Carlo simulation

# Model for chance nodes

- Expectiminimax search assumes to have a probabilistic model of how the opponent (or environment) will behave in any state
  - Model could be a simple uniform distribution (roll a die)
  - Model could be sophisticated and require a great deal of computation
  - We have a chance node for any outcome out of our control: opponent or environment
  - The model might say that adversarial actions are likely



# Games

- Classified over different axes:
  - Deterministic or stochastic
  - Perfect or imperfect information
  - Number of players
  - Zero sum

# Games with imperfect information

- Games could be partially observable
  - Result into belief states
- Solution for each state in the belief as deterministic game (with all techniques we have seen)
  - Instead of solving all of them, randomly samples states in the belief to solve them

# Commonsense - Information is important

Mondays:

- Road A leads to 1 gold. Road B leads to a fork.
- Go left from the fork, 100 gold. Go right, die.
- *Optimal strategy: B*


Tuesdays:

- Road A leads to 1 gold. Road B leads to a fork.
- Go left from the fork, die. Go right, 100 gold.
- *Optimal strategy: B*

I can't remember what day it is. *Choose B?*

# Multi-Agent Utilities

- With multiple players, the tree can be extended with additional nodes corresponding to each player
- Generalization of minimax:
  - Terminals have utility tuples
  - Node values are also utility tuples
  - Each player maximizes its own component
  - Can give rise to cooperation and competition dynamically



# Summary

- Stochastic games introduce chance nodes
  - Expectiminimax for finding the solution that maximizes the average case
  - A probabilistic model is needed
- Partially observable games result in belief states

# Next

- Non-zero sum games

# Games

- Classified over different axes:
  - Deterministic or stochastic
  - Perfect or imperfect information
  - Number of players
  - Zero sum

# Non-zero sum games

- In a Non-Zero-Sum Game, all parties could gain, or all parties could lose
  - means there's at least one outcome in which (A's PAYOFF + B's PAYOFF)  $\neq 0$
- Game theory to find a solution
  - Agent design: determine the best strategy against a rational player and the expected return for each player
- Real world applications: negotiations, bandwidth sharing, auctions, bankruptcy proceedings, pricing decisions



# Example: Prisoner's dilemma

- Alice and Bob have both been caught red handed near the scene of a burglary by the police and are interrogated separately
- The police present options:
  - Testify against the partner
  - Refuse to testify against the partner
- Consequences:
  - If you testify against your partner and your partner refuses, you are released and your partner will serve 10 years in jail
  - If you refuse and your partner testifies against you, you will serve 10 years in jail and your partner is released
  - If both of you testify against each other, both of you will serve 5 years in jail
  - If both of you refuse, both of you will only serve 1 year in jail
- No communication

# Prisoner's dilemma

- Payoff matrix to represent the utilities

	Bob: cooperate	Bob: defect
Alice: cooperate	A=-1, B=-1	A=-9, B=-0
Alice: defect	A=0, B=-9	A=-6, B=-6

What would you  
do if you were  
Alice?

# Prisoner's dilemma – Normal form

- Payoff matrix to represent the utilities

	Bob: cooperate	Bob: defect
Alice: cooperate	A=-1, B=-1	A=-9, B=0
Alice: defect	A=0, B=-9	A=-6, B=-6

$$n = 2$$

$$S_1 = \{C,D\}$$

$$S_2 = \{C,D\}$$

Each player has:

- Strategy space  $S_i$ , i.e., the available actions
- Payoff function, which given the strategies returns utility

$$u_1(C,C) = -1$$

$$u_1(C,D) = -9$$

$$u_1(D,C) = 0$$

$$u_1(D,D) = -6$$

$$u_2(C,C) = -1$$

$$u_2(C,D) = 0$$

$$u_2(D,C) = -9$$

$$u_2(D,D) = -6$$

# Normal Form Representation of a Non-Zero-Sum Game with $n$ players

Is a set of  $n$  strategy spaces  $S_1, S_2 \dots S_n$   
where  $S_i$  = The set of strategies available to player  $i$

And  $n$  payoff functions

$$u_1, u_2 \dots u_n$$

where

$$u_i : S_1 \times S_2 \times \dots \times S_n \rightarrow \mathbb{R}$$

is a function that takes a combination of strategies (one for each player) and returns the payoff for player  $i$

# Constraint Satisfaction

- Model problems as Constraint Satisfaction Problems
- Solve them through “backtracking” (depth-first search)

# Outline

- CSP model
- CSP search

# Outline

- CSP model
- CSP search

# Constraint satisfaction

- For some problems the “path” does not matter but the goal is important
- Real-world examples
  - Scheduling (job shops, section leader assignments, ...)
  - Placing component on chips
  - ...






# Example: Map coloring





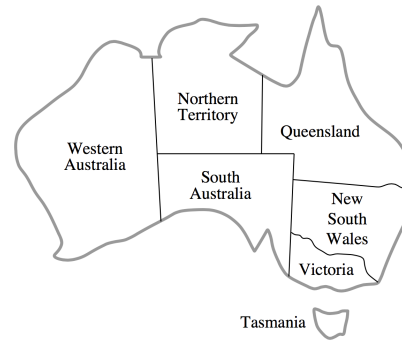
This is a royalty free image that can be used for your personal, corporate or education projects. It can not be resold or freely distributed, if you need an editable PowerPoint or Adobe Illustrator version of this map please visit [www.bjdesign.com](http://www.bjdesign.com) or [www.mapsfordesign.com](http://www.mapsfordesign.com). This text can be cropped off. © Copyright Bruce Jones Design Inc. 2010

# Example: Map coloring



# Example: Map coloring

- Variables: WA, NT, SA, Q, NSW, V, T
- Domains:  $D_i$ : r, g, b
- Constraints:
  - Explicit
    - $(WA, NT) \in \{(r, g), (r, b), (g, r), (g, b), (b, r), (b, g)\}$
    - $(NT, SA) \in \{(r, g), (r, b), (g, r), (g, b), (b, r), (b, g)\}$
    - ...
  - Implicit
    - $WA \neq NT$
    - $WA \neq SA$
    - ...
- Assignment:  
 $\{WA=r, NT=g, Q=r, NSW=g, V=r, SA=b, T=g\}$




# CSP definition

- $n$  variables:  $X_1, \dots, X_n$
- For each variable, a domain  $D_i$  of possible values.  
Example:  $D_1 : X_1 \in v_1, v_2, v_3$
- $m$  constraints  $C_1, \dots, C_n$ , each specifying allowable combinations of values for some set of variables
- An assignment of values to variables is the state of the problem.
  - We want to find a **complete** assignment (all variables with a value) **consistent** with the constraints

# Constraint graphs


- Binary CSP: each constraint relates at most two variables
- Nodes are variable, arcs show constraints



# Varieties of constraints


- Unary constraints involve a single variable (equivalent to reducing domains)
  - E.g., in the map coloring problem  $SA \neq g$
- Binary constraints involve pairs of variables, e.g.:
  - E.g., in the map coloring problem  $SA \neq WA$
- Higher-order constraints involve 3 or more variables (alldiff):
  - e.g., cryptarithmic column constraints
- Preferences (soft constraints):
  - E.g., red is better than green
  - Often representable by a cost for each variable assignment
  - Gives constrained optimization problems
  - We won't include them in the CSP

# Example: N-Queens



# Example: 4-Queens

- Formulation 1:
  - Variables:  $Q_1, Q_2, Q_3, Q_4$
  - Domains:  $Q_i = \{(x,y) \mid x \text{ in } [0,3] \text{ and } y \text{ in } [0,3]\}$
  - Constraints
    - Any queen not in the same x
      - $x_1 \neq x_2, x_1 \neq x_3, \dots$
    - Any queen not in the same y
      - $y_1 \neq y_2 \dots$
    - Any queen not on the diagonal
      - $\text{abs}(y_2 - y_1) \neq \text{abs}(x_2 - x_1) \dots$



Discussion



# Example: N-Queens

- Formulation 2:


- Variables:  $X_{ij}$
- Domains:  $\{0, 1\}$
- Constraints

$$\forall i, j, k \quad (X_{ij}, X_{ik}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{kj}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j+k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j-k}) \in \{(0, 0), (0, 1), (1, 0)\}$$



$$\sum_{i,j} X_{ij} = N$$

Discussion

# Example: N-Queens

- Formulation 3:

- Variables:  $Q_k$


- Domains:  $\{1, 2, 3, \dots, N\}$

- Constraints:

Implicit:  $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

Explicit:  $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$

...



# Example: Cryptarithmic

- Variables

$F \ T \ U \ W \ R \ O \ X_1 \ X_2 \ X_3$

- Domains:  $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$


- Constraints

$\text{alldiff}(F, T, U, W, R, O)$

$$\begin{aligned} O + O &= R + 10 \cdot C_{10} \\ C_{10} + W + W &= U + 10 \cdot C_{100} \\ C_{100} + T + T &= O + 10 \cdot C_{1000} \\ C_{1000} &= F, \end{aligned}$$


$$\begin{array}{r} \phantom{+} T \ W \ O \\ + \phantom{+} T \ W \ O \\ \hline F \ O \ U \ R \end{array}$$

$X_1$



Discussion

# Example: Sudoku



- Variables:
  - Each (open) square
- Domains:
  - $\{1,2,\dots,9\}$
- Constraints:

9-way alldiff for each column

9-way alldiff for each row

9-way alldiff for each region

(or can have a bunch of  
pairwise inequality  
constraints)

Discussion

# Example: Assignment problem

- Assign four workers W1,W2,W3,W4 to four products such that each worker works on one product and each product is produced by one worker.
- Effectiveness of production is given by the following table (e.g., worker W1 produces P1 with effectiveness 7) and the total effectiveness must be 19 at least

	P1	P2	P3	P4
W1	7	1	3	4
W2	8	2	5	1
W3	4	3	7	2
W4	3	1	6	3

# Example: assignment problem

- Variables
  - W1, W2, W3, W4
- Domains:
  - {P1, P2, P3, P4}
- Constraints
  - All\_diff(W1, W2, W3, W4)
  - $E1(W1) + E2(W2) + E3(W3) + E4(W4) \geq 19$

	P1	P2	P3	P4
W1	7	1	3	4
W2	8	2	5	1
W3	4	3	7	2
W4	3	1	6	3

# Example: The circuit-board layout problem

- Variables:
  - a variable for each piece
- Domains:
  - Possible placements over the board
- Constraints:
  - Needs to be in the board
  - Needs not to overlap with other pieces

